



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Parallel
Programming

Reminder

Moodle course

- If you are not enrolled in moodle, write an e-mail to Lukas
 - lukas.rothenberger@tu-darmstadt.de
 - <https://moodle.tu-darmstadt.de/course/view.php?id=40840>

Hand-ins

- Choose a group in moodle
- Put your matriculation number into moodle
- Hand in the signed forms for Lichtenberg
 - Not strictly required for course, but exercises will be graded on the system
 - Offers possibility of CUDA-compatible graphics card and many threads

Exercises

- Theoretical, in moodle:
 - On your own
 - Covers all material up to that week
- Practical labs:
 - In groups of three
- Oral tests (Testate):
 - On your own
 - Tutors might offer slots for groups nevertheless

Consultation hours

- Find the consultation hours of the tutors in moodle:
 - Allgemeines > Sprechstunden
- Tutors can change location during the semester



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Parallel
Programming

Programming in C++

C++ is the upgrade

- Bjarne Stroustrup started in 1979 to extend C with classes
- Standardized in 1998 and 2003
- Since 2011 standardized every three years
 - Most recent standard: C++23
 - C++26 is a work in progress

Where to look

<https://en.cppreference.com>

Features of C++

- You cannot solve a complex problem with simple tools
- We will highlight multiple interesting features
 - Most advanced or dangerous topics avoided

Philosophy

- “Zero-overhead principle”
 - You do not have to pay for what you do not use
 - If you were to program a feature by hand, you would not achieve better
- A ‘large’ standard
 - What is impossible to achieve on your own
 - What is hard to get right
 - What is popular

Warning

We do not follow (our own) code-style guidelines!

References

The drawbacks of pointers

- Pointers can be uninitialized
- Pointers can point to garbage (incremented too many times)
- Pointers allow type shenanigans

The idea of a reference

- Have different names for the same memory cell
 - “Lukas”
 - “Herr Rothenberger”
 - “You” (furiously point finger)
- Think of it as a tool/way to access the memory cell

Let's start simple

```
int return_two(int number) {  
    int& name = number;  
  
    name = 2;  
  
    return number;  
}
```

type& is the reference for type

name now refers to number

name does not have an own memory location!

References preserve `const`

```
int return_two_failure(int number) {  
    const int& name = number;  
  
    name = 2;  
  
    int const& nickname = number;  
    nickname = 2;  
  
    return number;  
}
```

`const int&` and
`int const&` are similar

Cannot change the value via
name or nickname

The benefits of references

- They are always initialized
- Under the hood, they are an address
 - No hefty copy



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Parallel
Programming

Throws

The drawbacks or return values

- When we encounter a bad state in C, we need to propagate it via the return value of the function
- You always had to check for the return values
 - Obfuscated actual meaning of function

The idea of exceptions

- For exceptional behavior, we should not clutter the control path
- Distinguish return value from error code

Let's start simple

```
std::uint16_t convert(int number) {  
    if (number < 0) {  
        throw -1;  
    }  
    if (number > 65535) {  
        throw "Too high";  
    }  
  
    return std::uint16_t{number};  
}
```

We can **throw** anything

catch the culprit

```
std::uint16_t maybe() {  
    try {  
        return convert(-12);  
    } catch (const char *c) {  
        return 12;  
    } catch (int& i) {  
        return 13;  
    }  
}
```

If an expression in a **try** block throws, search for the suitable **catch**

We can **catch** anything

If there is no suitable **catch**, we keep **throw-ing**

There is also **finally**

The benefit of throwing

- If it is not your responsibility, you can ignore it
 - System out of memory?
 - Passed illegal argument?

Templates

The drawback of `void*`

- In C, we were not able to handle different types in the same function
 - We can take an argument as `void*`, but then we had to cast it manually
 - Some performance issues with `void*`

The idea of a template

- We provide a prototype for a function (later class) with some ‘blank spaces’ to be filled in later
- We specify what we need from the blank spaces

Let's start simple

```
template<typename _Ty>  
std::size_t get_size(_Ty object) {  
    std::size_t s = sizeof(object)  
    return s;  
}
```

We provide this function for
every type

The compiler will generate
an instance for each
invocation

Calling templated functions

```
void print_sizes() {  
    std::size_t s1 = get_size<int>(32);  
    std::size_t s2 = get_size<std::int64_t>(32);  
    std::size_t s3 = get_size(32);  
  
    ... (printing comes later)  
}
```

Restrictions on templates

```
template<typename _Ty>
_Ty increment_to_one() {
    _Ty zero = _Ty(0);
    zero++;
    return zero;
}
```

We require that there is a constructor of `_Ty` taking an `int`

We require that we can call `++` on an object of type `_Ty`

We could also call any method imaginable

Value templates

- Our ‘blank spaces’ were types up until now
- But there is not much hindering as to template on values
 - Other than rounding

Taking a value

```
template<int _val, typename _Ty>
int increment(_Ty value) {
    value += _val;
    return value;
}
```

More work during compile
time

The benefits of templates

- Now, we can write containers, adaptable algorithms, etc.
- We catch type errors during compile time



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Parallel
Programming

Classes



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Parallel
Programming

and resources



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Parallel
Programming

and operators

The drawbacks of structures

- Disconnect between data and functionality
- Dependencies between data not expressible
- No customization of operators

The drawbacks of resources

- You always have to keep book about the resources you opened
- What's the convention if you call other libraries?

The idea of classes

- Group data and functionality together
- Hide unnecessary information from others
- (Allow polymorphism, dynamic dispatch, etc.)
- Have fancy syntax

Let's start 'simple'

```
template<typename _Ty1, typename _Ty2>  
class pair {  
  
_Ty1 val1;  
  
_Ty2 val2;  
  
  
public:  
pair(const _Ty1& v1, const _Ty2& v2)  
    : val1{ v1 }, val2{ v2 }  
    {}  
};
```

val1 and val2 are
private members

Note the member variable
initialization

Remember this class

```
template<typename _Ty1, typename _Ty2>  
class pair {  
    _Ty1 val1;  
    _Ty2 val2;  
  
public:  
    pair(const _Ty1& v1, const _Ty2& v2)  
        : val1{ v1 }, val2{ v2 }  
    {}  
};
```

Declaring methods

```
...  
  
_Ty1& get_first_ref() {  
  
    return val1;  
  
}  
  
  
const _Ty1& get_first_ref() const {  
  
    return val1;  
  
}  
  
...
```

The compiler chooses the method based on the object

Differing implementations possible

Fun with operators

- In C++, you can define operators for classes
 - Well known: `=`, `==`, `<=`, `++`, `--`, `[]`
 - Curious: casting, `&`, `(...)`, `,`

Copying

- What must I do to copy an object of my class?
- The default way is to recursively copy member-wise
 - C-types are copied trivially

Copying resources

```
template<typename _Ty1, std::size_t _si>
class my_array {
    _Ty1 *pointer;

public:
    my_array() : pointer{ new _Ty1[_si] }
    {
    };
}
```

`new` is C++'s way of `malloc`
– it is typed!

A trivial copy fails

Let's fail

```
int main(int argc, char **argv) {  
    my_array<int, 10> arr{};  
    my_array<int, 10> copy = arr;  
    if (argc < 2) {  
        my_array<int, 10> other{};  
        copy = other;  
    }  
  
    return 3;  
}
```

We copied shallowly

We should ask in the forum!

Remember this class

```
template<typename _Ty, std::size_t _si>
class my_array {
    _Ty *pointer;

public:
    my_array() : pointer{ new _Ty[_si] }
    {
    };
}
```

Let's construct successfully

```
...  
my_array(const my_array& other)  
    : pointer{ new _Ty1[_si] } {  
        for (std::size_t i = 0; i < _si; i++) {  
            pointer[i] = other.pointer[i];  
        }  
    }  
...  
}
```

We take other by constant reference

If we were to take it by value, we would need to construct other first...

State of the union

```
int main(int argc, char **argv) {  
    my_array<int, 10> arr{};  
  
    my_array<int, 10> copy = arr;  
  
    if (argc < 2) {  
  
        my_array<int, 10> other{};  
  
        copy = other;  
  
    }  
  
    return 2;  
}
```

We have correctly
constructed copy

Let's copy successfully

```
...  
my_array& operator=(const my_array& other) {  
    for (std::size_t i = 0; i < _si; i++) {  
        pointer[i] = other.pointer[i];  
    }  
    return *this;  
}  
...
```

`this` is always a pointer to
the class

`*this` is a reference

Values are a lie

State of the union

```
int main(int argc, char **argv) {  
    my_array<int, 10> arr{};  
  
    my_array<int, 10> copy = arr;  
  
    if (argc < 2) {  
        my_array<int, 10> other{};  
  
        copy = other;  
  
    }  
  
    return 1;  
}
```

We have correctly
constructed copy

We have correctly
assigned copy

Returning memory

- Remember your responsibility from C?
 - If you ask for memory yourself, you are responsible
 - `malloc()` needed a call to `free()`
- We asked for memory with `new[]`
 - We now need to return it with `delete[]`
- Do not mix and match!
 - There is also the pair `new` and `delete`

Let's destruct successfully

```
...  
~my_array() {  
    delete[] pointer;  
}  
...
```

The runtime knows the size
of the memory

The destructor

- There can be many ways to construct an object
 - Different overloads of the constructor allowed
- There always can only be one way to destruct an object
 - No arguments
 - Called immediately when it goes out of scope
 - Clean up all you are responsible for

State of the union

```
int main(int argc, char **argv) {  
    my_array<int, 10> arr{};  
  
    my_array<int, 10> copy = arr;  
  
    if (argc < 2) {  
        my_array<int, 10> other{};  
  
        copy = other;  
    }  
  
    return 0;  
}
```

We have correctly
constructed copy

We have correctly
assigned copy

We have correctly
cleaned up

Accessing our data

```
...  
  
_Ty operator[](std::size_t index) {  
  
    if (index >= _si) {  
  
        throw std::exception{};  
  
    }  
  
    return pointer[index];  
  
}  
  
...
```

Some operators have fixed parameters, some not

C++23: Allows multiple std::size_t

Before: cheating possible

The benefit of classes

- A whole new world opens up!
 - Automatic resource management
 - Great code quality

Various smaller topics

Standard containers

The all 'rounder': `std::vector<_Ty>`

- Managed block of contiguous memory for objects of type `_Ty`
- Constant access time
- Amortized constant time for pushing elements at the end
- Constant time for deleting elements at the end

The all ‘ugly: `std::list<_Ty>`

- Linked list of nodes containing one object each
- Linear access time
- Constant merging
- Nearly always slower than `std::vector<_Ty>`

The fixer: `std::array<_Ty, _si>`

- You get exactly what you expect

The bouncer: `std::queue<_Ty>`

- First in, first out
- Nothing more

The dishwasher: `std::stack<_Ty>`

- First in, last out
- Nothing more

The odd uncle: `std::deque<_Ty>`

- List of contiguous memory blocks
- Constant access time
- Constant push and pop at both ends
- No one really knows what's going on

The matchmaker: `std::map<_Ke, _Val>`

- In other programming languages known as a dictionary
- Associative container mapping `_Ke` to `_Val`
 - Uses `operator[]`
- Logarithmic access time, stored as red–black tree

The better matchmaker: `std::unordered_map<_Key, _Value>`

- In other programming languages known as a hash map
- Amortized constant insert/access/removal time

The savior: `std::string`

- Yeah, there is a string type in C++
- It offers a lot of functionality

Iterating a container

- Most containers offer:
 - `begin()`, `begin() const`, `cbegin() const`
 - `end()`, `end() const`, `cend() const`
- They return a (constant) iterator
 - Proxy object to access the container
 - Point to the first/one-after-last element
 - Can increment iterator with `++`
 - Can retrieve element with `*`

Let's start 'simple'

```
template<typename _Contain>
double sum(const _Contain& container) {
    auto temp = 0.0;
    for (auto it = container.begin();
          it != container.end(); ++it) {
        temp += *it;
    }
    return temp;
}
```

`auto` means:
Compiler, do your thing!

We require a lot of `_Contain`

auto in action

- `auto` matches greedily any type
 - Might want to use `const auto` or `const auto&`
- `auto` is necessary for types that cannot be named

Let's get complex

```
template<typename _Ke>
double sum(const std::map<_Ke, double>& m) {
    auto temp = 0.0;
    for (const auto& [_, val] : m) {
        temp += val;
    }
    return temp;
}
```

for (type name : c)
is C++'s foreach

auto [a, b] 'matches'
trivial pairs

auto [a, b, c] 'matches'
trivial triples

The benefits of standard container

- They are really hard to get right!
- Tuned by many smart minds
 - Small-vector optimization for some libraries
 - Exception-safe

I/O

Template + I/O

- Methods like `printf` and `scanf` are kind-of incompatible with templates
 - What specifier would you insert into the formatting string?

C++'s I/O

```
template<typename _Ty>
void print(const _Ty& val) {
    std::cout << val << '\n';
    _Ty copy {};
    std::cin >> copy;
    std::cout << (val == copy) << std::endl;
}
```

<< and >> shove data into
our out of the stream

Can overload std::cout
and std::cin for custom
types



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Parallel
Programming

Dependent types

What is necessary?

- C++ goes hard when it comes to types
 - You have to convert between types all the time
 - In C, all had the same type (and you had to care for the values)

std::vector<T>

Member types

Member type	Definition
<code>value_type</code>	<code>T</code>
<code>allocator_type</code>	<code>Allocator</code>
<code>size_type</code>	Unsigned integer type (usually <code>std::size_t</code>)
<code>difference_type</code>	Signed integer type (usually <code>std::ptrdiff_t</code>)
<code>reference</code>	<code>value_type&</code>
<code>const_reference</code>	<code>const value_type&</code>
<code>pointer</code>	<code>Allocator::pointer</code> (until C++11) <code>std::allocator_traits<Allocator>::pointer</code> (since C++11)
<code>const_pointer</code>	<code>Allocator::const_pointer</code> (until C++11) <code>std::allocator_traits<Allocator>::const_pointer</code> (since C++11)
<code>iterator</code>	<code>LegacyRandomAccessIterator</code> and <code>LegacyContiguousIterator</code> to <code>value_type</code> (until C++20) <code>LegacyRandomAccessIterator</code> , <code>contiguous_iterator</code> , and <code>ConstexprIterator</code> to <code>value_type</code> (since C++20)
<code>const_iterator</code>	<code>LegacyRandomAccessIterator</code> and <code>LegacyContiguousIterator</code> to <code>const value_type</code> (until C++20) <code>LegacyRandomAccessIterator</code> , <code>contiguous_iterator</code> , and <code>ConstexprIterator</code> to <code>const value_type</code> (since C++20)
<code>reverse_iterator</code>	<code>std::reverse_iterator<iterator></code>
<code>const_reverse_iterator</code>	<code>std::reverse_iterator<const_iterator></code>

The index types

- All indexed containers offer `::size_type`
 - Usually `std::size_t`
 - `std::size_t` is large enough to index any contiguous memory
 - `std::vector<HUGE_TYPE>::size_type` can be smaller



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Parallel
Programming

The ‘correct’ method

C++ is opt-in

```
class holder {  
public:  
    [[nodiscard]] constexpr double get_value(  
        const int argument) const noexcept {  
        return 0.0;  
    }  
}
```

`[[nodiscard]]`: Warning if value not saved when called

`constexpr`: Can be evaluated during compile time

`noexcept`: Does not throw any value

Lambdas

Purpose of lambdas

- Sometimes, you need a small function
 - Not necessarily a name
 - Usually not passed around

Let's start simple

```
int main() {  
    const std::vector<int> v {1, 2, 3, 4, 5};  
    const int mod_value = 3;  
  
    const auto pos = std::find_if(  
        v.begin(), v.end(),  
        [mod_value] (const int x)  
        {return x % mod_value == 0;});  
  
    return pos != v.end();  
}x
```

[...] contains the values from outside that are made available. & possible

(...) contains the arguments

{...} contains the instructions



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Parallel
Programming

A case study

Our setting

- We have a directed graph $G = (V, A)$ whose nodes have some position in Euclidean space
- We want to calculate the average Euclidean length of the arcs
- The graph and the calculation are distributed with MPI
 - As MPI is not part of this year's syllabus, we accept the functions as black boxes

The class for positions : 1

```
template <typename T>
    requires std::is_arithmetic_v<T>
class Vec3 {
    T x{0};
    T y{0};
    T z{0};

public:
    using value_type = T;

    constexpr Vec3() = default;
    constexpr ~Vec3() = default;
```

We do not accept any type,
it must be arithmetic

We initialize the members to
0, no matter the constructor
call

We provide the dependent
type `value_type`

We want default
construction/destruction of
our class

The class for positions : 2

```

constexpr explicit Vec3(const T& val) noexcept
|   : x(val),
|   y(val),
|   z(val) {
|
}

constexpr Vec3(const T& _x, const T& _y, const T& _z)
|   noexcept : x(_x), y(_y), z(_z) { }

constexpr Vec3(const Vec3<T>& other)
= default;
constexpr Vec3<T>& operator=(const Vec3<T>& other)
= default;

```

We allow construction by one or three values

We do not want implicit conversion, so we use **explicit**

We have no resources in the class, so we use **default**

The class for positions : 3

```
[[nodiscard]] constexpr friend
    Vec3<T> operator-
    (const Vec3<T>& lhs, const Vec3<T>& rhs)
noexcept {
    return {
        lhs.x - rhs.x,
        lhs.y - rhs.y,
        lhs.z - rhs.z
    };
}
```

We want to subtract vectors,
friend makes the operator
access the private members

return {...} can often
deduce the constructor

The class for positions : 4

```
[[nodiscard]] double calculate_2_norm() const noexcept {
    return std::sqrt(calculate_squared_2_norm());
}
```

Sometimes, you want to store intermediate values for debuggability

```
[[nodiscard]] constexpr double calculate_squared_2_norm()
    const noexcept {
    const auto xx = x * x;
    const auto yy = y * y;
    const auto zz = z * z;

    const auto sum = xx + yy + zz;
    return sum;
}
```

Let the optimizer do the heavy-lifting

std::sqrt is not `constexpr`

Calculate distance for one node

```
[[nodiscard]] static double compute_arc_length
(const DistributedGraph& graph, const node_id_type node_id) {
    const auto my_rank = MPIWrapper::get_my_rank();

    auto accumulated_distance = 0.0;
    const auto& node_position = graph.get_node_position(my_rank, node_id);
    const auto out_arcs = graph.get_out_arcs(my_rank, node_id);

    for (const auto& [target_rank, target_id, weight] : out_arcs) {
        const auto& position = graph.get_node_position(target_rank, target_id);

        const auto difference = position - node_position;
        const auto distance = difference.calculate_2_norm();

        accumulated_distance += distance;
    }

    return accumulated_distance;
}
```

get_node_position()
returns Vec3<double>

get_out_arcs() returns a
std::span

std::span functions as a
view on contiguous data

Calculating distance for all

```
[[nodiscard]] static double all_compute_arc_length(const DistributedGraph& graph) {
    const auto number_local_nodes = graph.get_number_local_nodes();
    auto accumulated_distance = 0.0;

    for (auto node_id = node_id_type{0}; node_id < number_local_nodes; ++node_id) {
        const auto distance = compute_arc_length(graph, node_id);
        accumulated_distance += distance;
    }

    const auto total_distance = MPIWrapper::all_reduce_sum(accumulated_distance);
    const auto number_total_arcs = OutArcCounter::all_count_out_arcs(graph);

    const auto& average_distance = total_distance / number_total_arcs;

    return average_distance;
}
```

all_reduce_sum() takes local values and returns the sum

all_count_out_arcs() returns the number of arcs in the graph

Takeaways

- Let the compiler do the heavy lifting
- Name your functions/methods/variables with **common sense**
- Break your problem into suitable sub-problems